

Optimisation du Game of Life sur GPU

De l'implémentation naïve au Bit-Packing

A5 TP Conway

NVIDIA RTX 3080 — Grille 32k × 32k

48 FPS → 416 FPS

Version 1: noif-char (Naïve)

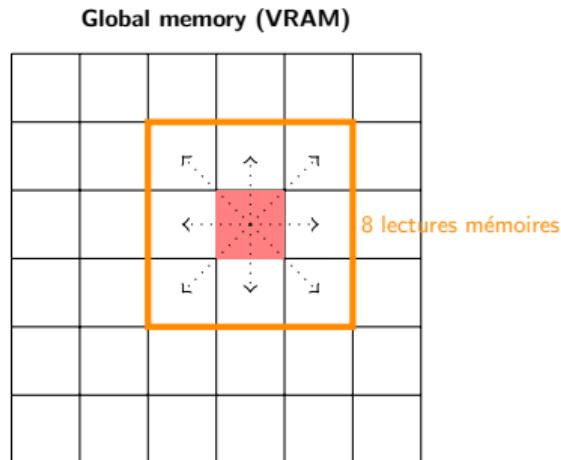
Taille grille: 32768×32768 octets
 ≈ 1 Go de RAM

FPS: 48.69

Type des données: char (1 octet = 1 cellule)

Pas de shared memory

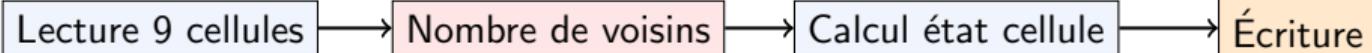
Taille bloc: 32×32 threads (calculent 32×32 cellules)



Logique “branchless” (noif)

```
// On compte les voisins vivants
int neighbors = 0;
for (int dy = -1; dy < 2; dy++) {
    for (int dx = -1; dx < 2; dx++) {
        if (x + dx >= 0 && x + dx < width && y + dy >= 0 && y + dy < height && (dy != 0 || dx != 0)) {
            neighbors += grid[idx + dx + dy * width];
        }
    }
}

// Calcul conditionnel sans 'if'
int alive = (neighbors == 3) | (grid[y * width + x] == 1 && neighbors == 2);
new_grid[idx] = alive;
```

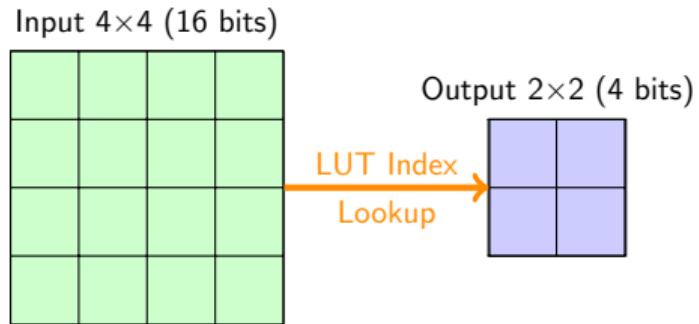


Version 2: Lookup table 4×4

Performance:

- **FPS:** 112.76 ($\times 2.3$)
- **Stratégie:** Pré-calcul (LUT)
- **Mémoire:** Utilisation de la **shared memory**

Concept: Réduire les calculs (comptage voisins) par une simple lecture tableau.



```
unsigned char host_lut[65536]
```

Construction de l'index LUT

```
// Construction de l'index 16-bits
uint16_t state_idx = 0;
#pragma unroll
for (int r = 0; r < 4; r++) {
    #pragma unroll
    for (int c = 0; c < 4; c++) {
        if (tile[tile_r + r][tile_c + c]) {
            state_idx |= (1 << (r * 4 + c));
        }
    }
}
unsigned char res = lut[state_idx];
```

| | | | |
|-----|-----|-----|-----|
| b30 | b31 | b32 | b33 |
| b20 | b21 | b22 | b23 |
| b10 | b11 | b12 | b13 |
| b00 | b01 | b02 | b03 |

→ uint16_t state_idx

Bloc 4×4 (voisinage)

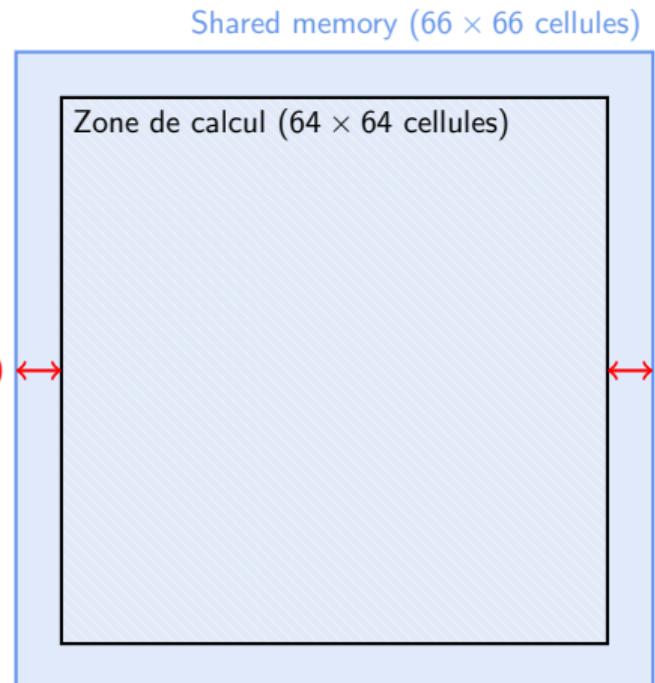
Gestion mémoire: stratégie par bloc

Un thread calcule **2×2 cellules**

Bloc 32×32 threads

Surface calculée: **64×64 cellules**

Shared Memory: **66×66 cellules**



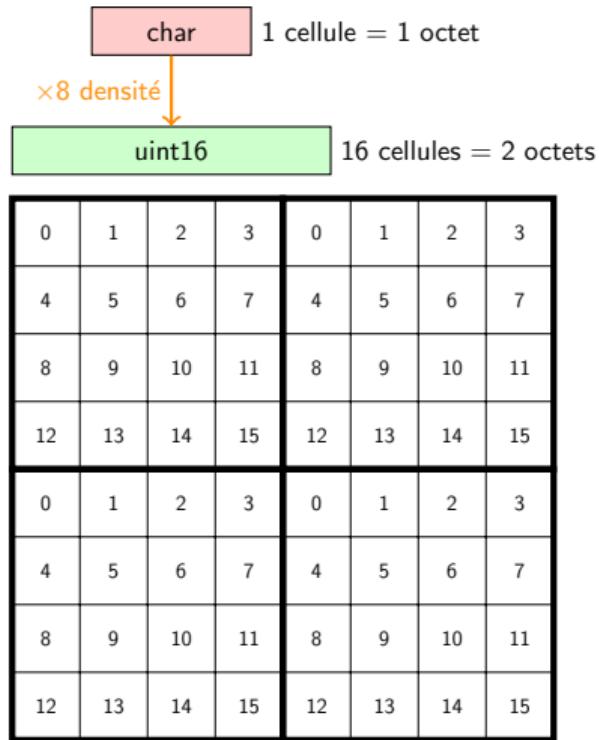
Version 3: uint16 (Bit-Packing)

Performance:

- **FPS: 416.51** ($\times 8.5$ vs naïve)
- **Stratégie:** compression mémoire

Principe:

- 1 cellule = 1 bit (pas 1 octet)
- Grille 32768×32768 cellules
- $\frac{32768 \times 32768}{8 \text{ bits/octet}} = 128 \text{ Mo}$



Architecture du kernel uint16

1 uint16 = 1 thread

| | | | |
|----|----|----|----|
| 0 | 1 | 2 | 3 |
| 4 | 5 | 6 | 7 |
| 8 | 9 | 10 | 11 |
| 12 | 13 | 14 | 15 |

Architecture du kernel uint16

1 uint16 = 1 thread

Utilisation de la LUT 4×4

| | | | |
|----|----|----|----|
| 0 | 1 | 2 | 3 |
| 4 | 5 | 6 | 7 |
| 8 | 9 | 10 | 11 |
| 12 | 13 | 14 | 15 |

Architecture du kernel uint16

Utilisation de la LUT 4×4
Construction de l'indice 16-bits
à partir des uint16 voisins

1 uint16 = 1 thread

| | | | |
|----|----|----|----|
| 0 | 1 | 2 | 3 |
| 4 | 5 | 6 | 7 |
| 8 | 9 | 10 | 11 |
| 12 | 13 | 14 | 15 |

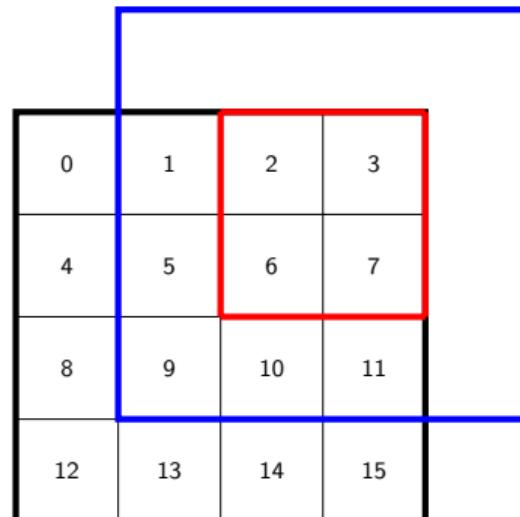
Architecture du kernel uint16

Utilisation de la LUT 4×4

Construction de l'indice 16-bits
à partir des uint16 voisins

Patch suivant

$1 \text{ uint16} = 1 \text{ thread}$



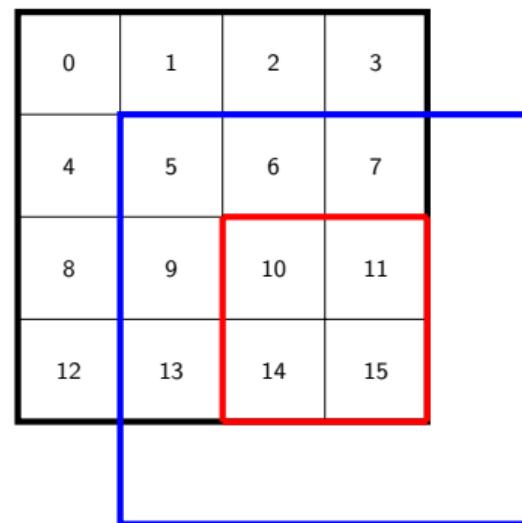
Architecture du kernel uint16

1 uint16 = 1 thread

Utilisation de la LUT 4×4

Construction de l'indice 16-bits
à partir des uint16 voisins

Patch suivant



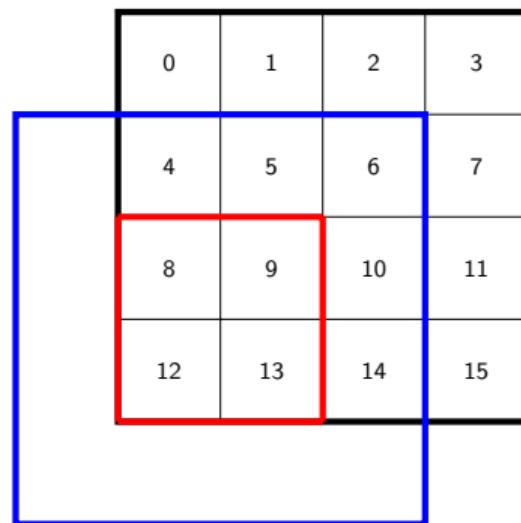
Architecture du kernel uint16

1 uint16 = 1 thread

Utilisation de la LUT 4×4

Construction de l'indice 16-bits
à partir des uint16 voisins

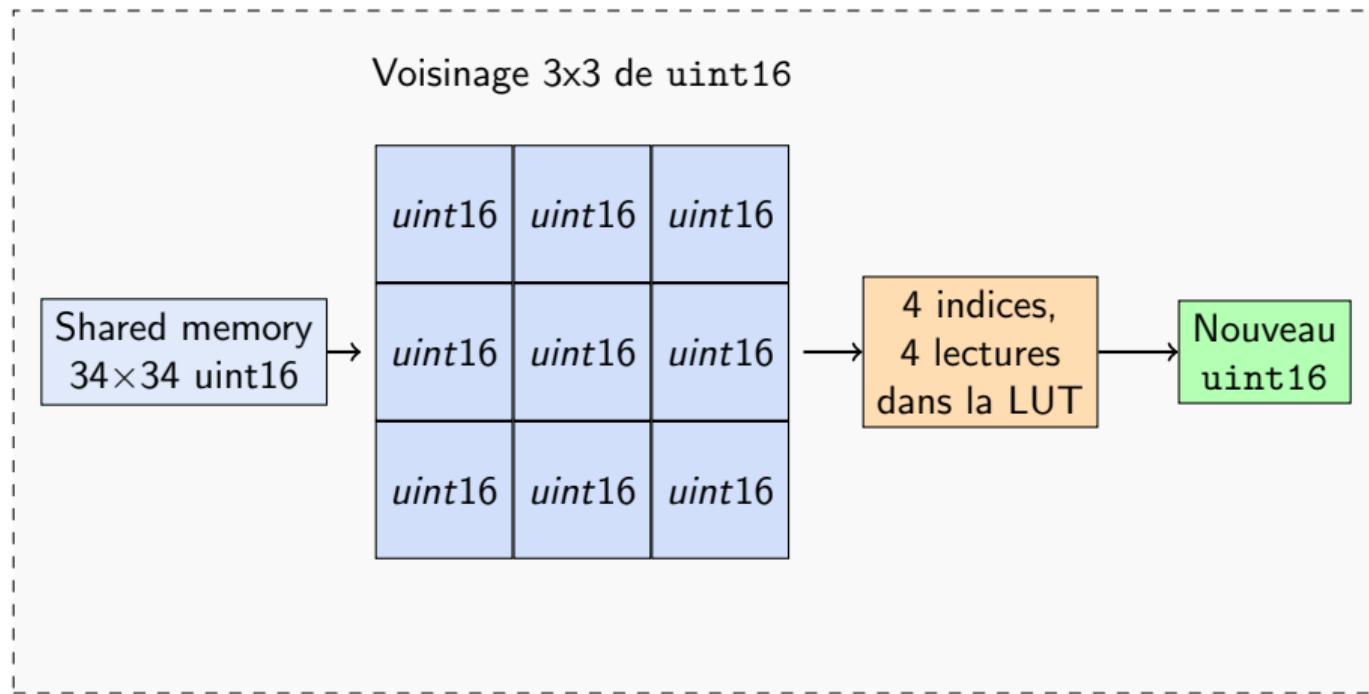
Patch suivant



Architecture du kernel uint16

1 thread = uint16

Bloc 32×32 threads = 128×128 cellules



Résumé des gains de performance

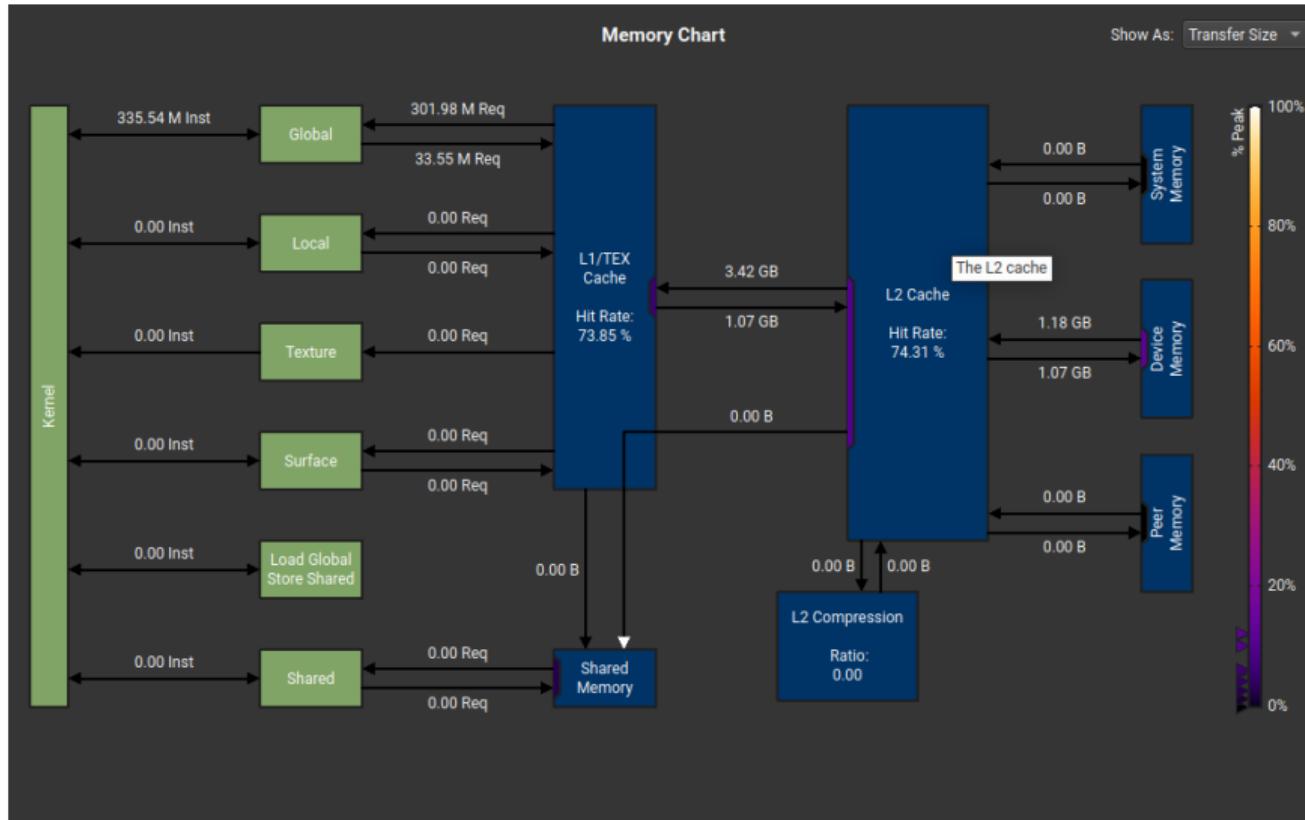
Métriques de débit et temps d'exécution

| Version | Thread/cellule | Durée (ms) | BP Mém. (Go/s) | BP Mém. (%) | FPS |
|-----------|----------------|-------------|----------------|-------------|--------|
| noif-char | 1 | 25.21 | 89.42 | 32.72 | 48.81 |
| lookup4×4 | 4 | 11.21 | 196.02 | 48.30 | 152.04 |
| uint16 | 16 | 3.03 | 90.54 | 25.85 | 416.51 |

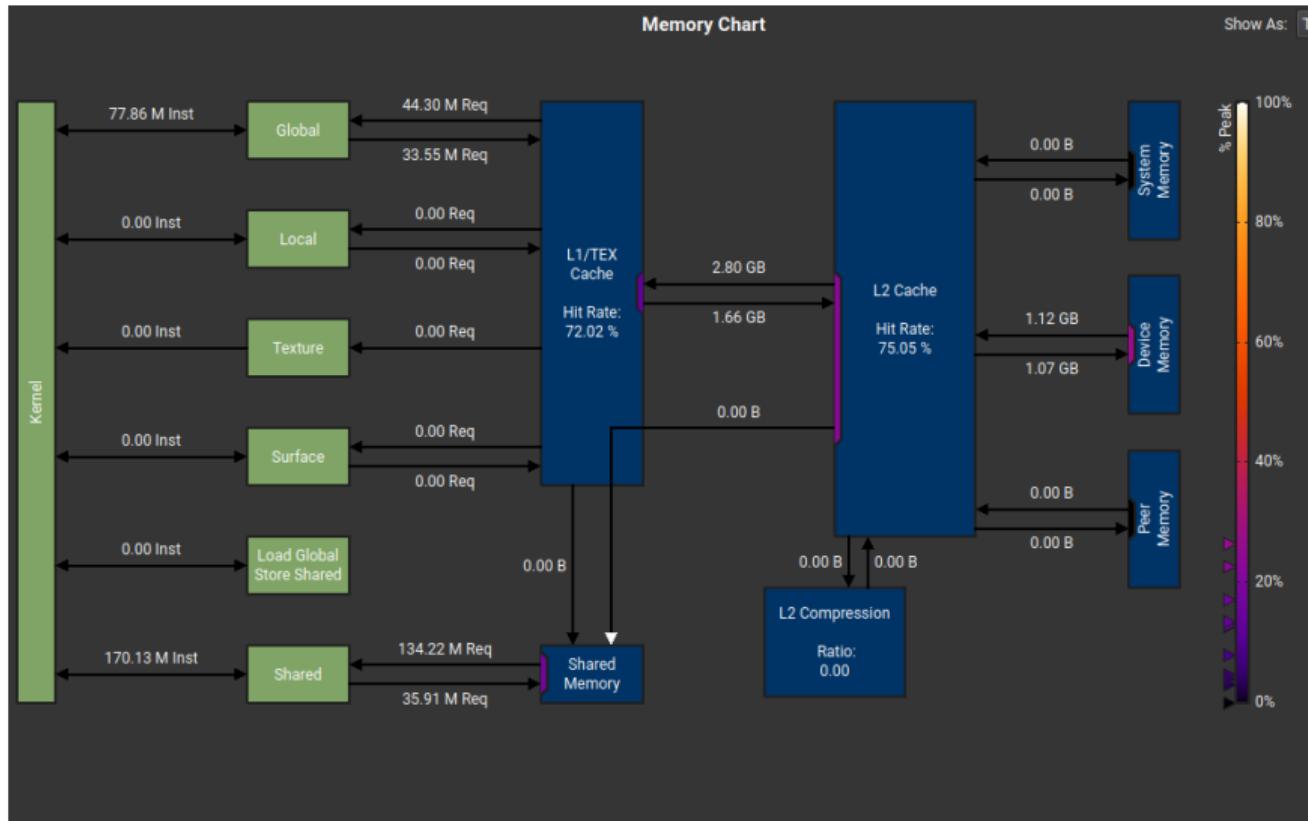
Efficacité matérielle

| Version | Hit L1/TEX (%) | Occupation Mém. (%) | SM Busy (%) | Calcul (%) |
|-----------|----------------|---------------------|-------------|------------|
| noif-char | 73.85 | 16.56 | 32.46 | 32.72 |
| lookup4×4 | 72.02 | 31.05 | 46.56 | 48.30 |
| uint16 | 95.56 | 20.68 | 39.18 | 37.25 |

Profiling: noif-char



Profiling: lookup4×4



Profiling: uint16 (division par 8 des accès mémoire)

